

```
In [2]: import numpy as np
import pandas as pd
```

```
In [3]: df=pd.read_csv("C:/Users/DELL/Desktop/india.csv")
```

```
In [4]: df.head()
```

Out[4]:

	HSCode	Commodity	value	country	year
0	5	PRODUCTS OF ANIMAL ORIGIN, NOT ELSEWHERE SPECI...	0.00	AFGHANISTAN TIS	2018
1	7	EDIBLE VEGETABLES AND CERTAIN ROOTS AND TUBERS.	12.38	AFGHANISTAN TIS	2018
2	8	EDIBLE FRUIT AND NUTS; PEEL OR CITRUS FRUIT OR...	268.60	AFGHANISTAN TIS	2018
3	9	COFFEE, TEA, MATE AND SPICES.	35.48	AFGHANISTAN TIS	2018
4	11	PRODUCTS OF THE MILLING INDUSTRY; MALT; STARCH...	NaN	AFGHANISTAN TIS	2018

```
In [5]: df_cat=df.select_dtypes(include=[object])
df_num=df.select_dtypes(include=[np.number])
```

```
In [6]: df_cat.head()
```

Out[6]:

	Commodity	country
0	PRODUCTS OF ANIMAL ORIGIN, NOT ELSEWHERE SPECI...	AFGHANISTAN TIS
1	EDIBLE VEGETABLES AND CERTAIN ROOTS AND TUBERS.	AFGHANISTAN TIS
2	EDIBLE FRUIT AND NUTS; PEEL OR CITRUS FRUIT OR...	AFGHANISTAN TIS
3	COFFEE, TEA, MATE AND SPICES.	AFGHANISTAN TIS
4	PRODUCTS OF THE MILLING INDUSTRY; MALT; STARCH...	AFGHANISTAN TIS

```
In [7]: df_num.head()
```

Out[7]:

	HSCode	value	year
0	5	0.00	2018
1	7	12.38	2018
2	8	268.60	2018
3	9	35.48	2018
4	11	NaN	2018

```
In [8]: df_cat.isnull().sum()
```

```
Out[8]: Commodity    0  
country      0  
dtype: int64
```

```
In [9]: df_num.isnull().sum()
```

```
Out[9]: HSCode      0  
value    14027  
year      0  
dtype: int64
```

```
In [10]: df_num.value.fillna(df_num.value.mean(),inplace=True)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\generic.py:6130: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
self._update_inplace(new_data)

```
In [11]: df_num.head()
```

```
Out[11]:
```

	HSCode	value	year
0	5	0.000000	2018
1	7	12.380000	2018
2	8	268.600000	2018
3	9	35.480000	2018
4	11	63.289855	2018

```
In [12]: from sklearn.preprocessing import LabelEncoder
```

```
In [13]: le=LabelEncoder()
```

```
In [14]: df_cat=df_cat.apply(le.fit_transform)
```

```
In [15]: df=pd.concat([df_cat,df_num],axis=1)
```

In [16]: `df.head()`

Out[16]:

	Commodity	country	HSCode	value	year
0	71	0	5	0.000000	2018
1	22	0	7	12.380000	2018
2	21	0	8	268.600000	2018
3	16	0	9	35.480000	2018
4	72	0	11	63.289855	2018

In [17]: `from sklearn.linear_model import LogisticRegression`

In [18]: `lr=LogisticRegression()`

In [19]: `x=df.iloc[:,[0,1,2,3,4]].values`
`y=df.iloc[:,2].values`

In [20]: `from sklearn.model_selection import train_test_split`

In [21]: `x_train,x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random_state=0) #shift+tab`

In [22]: `lr.fit(x_train,y_train)`

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:460: FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify the multi_class option to silence this warning.

"this warning.", FutureWarning)

Out[22]: `LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True, intercept_scaling=1, max_iter=100, multi_class='warn', n_jobs=None, penalty='l2', random_state=None, solver='warn', tol=0.0001, verbose=0, warm_start=False)`

In [23]: `pre=lr.predict(x_test)`

In [24]: `pre`

Out[24]: `array([42, 17, 17, ..., 84, 39, 72], dtype=int64)`

In [25]: `from sklearn.metrics import accuracy_score`

In [26]: `a=accuracy_score(pre,y_test)`

In [27]: a

Out[27]: 0.4576199466180587

In [28]: fit=lr.fit(x,y)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:460: FutureWarning: Default multi_class will be changed to 'auto' in 0.22. Specify the multi_class option to silence this warning.

"this warning.", FutureWarning)

In [29]: fit.score(x,y)

Out[29]: 0.4674579730382942

decision tree

In [30]: df.head()

Out[30]:

	Commodity	country	HSCode	value	year
0	71	0	5	0.000000	2018
1	22	0	7	12.380000	2018
2	21	0	8	268.600000	2018
3	16	0	9	35.480000	2018
4	72	0	11	63.289855	2018

In [31]: x1=df.iloc[:,[0,1,2,3,4]].values
y=df.iloc[:,2].values

In [32]: from sklearn.model_selection import train_test_split

In [33]: x1_train, x1_test, y_train, y_test = train_test_split(x1, y, test_size=0.33, random_state=0)

In [34]: from sklearn.tree import DecisionTreeClassifier

In [35]: dt=DecisionTreeClassifier()

```
In [36]: dt.fit(x1_train,y_train)
```

```
Out[36]: DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=None,
                                max_features=None, max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, presort=False, random_state=None,
                                splitter='best')
```

```
In [37]: pre=dt.predict(x1_test)
```

```
In [38]: pre
```

```
Out[38]: array([42, 17, 17, ..., 75, 49, 54], dtype=int64)
```

```
In [39]: from sklearn.metrics import accuracy_score
```

```
In [40]: b=accuracy_score(y_test,pre)
```

```
In [42]: b
```

```
Out[42]: 1.0
```

random forest

```
In [43]: df.head()
```

```
Out[43]:
```

	Commodity	country	HSCode	value	year
0	71	0	5	0.000000	2018
1	22	0	7	12.380000	2018
2	21	0	8	268.600000	2018
3	16	0	9	35.480000	2018
4	72	0	11	63.289855	2018

```
In [44]: x2=df.iloc[:,[0,1,2,3,4]].values
         y=df.iloc[:,2].values
```

```
In [45]: from sklearn.model_selection import train_test_split
```

```
In [46]: x2_train, x2_test, y_train, y_test = train_test_split(x2, y, test_size=0.33, r
                    randon_state=0)
```

```
In [47]: from sklearn.ensemble import RandomForestClassifier
```

```
In [48]: rf=RandomForestClassifier()
```

```
In [49]: rf.fit(x2_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\ensemble\forest.py:246: FutureWarning: The default value of n_estimators will change from 10 in version 0.20 to 100 in 0.22.

"10 in version 0.20 to 100 in 0.22.", FutureWarning)

```
Out[49]: RandomForestClassifier(bootstrap=True, class_weight=None, criterion='gini',
                                max_depth=None, max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=None,
                                oob_score=False, random_state=None, verbose=0,
                                warm_start=False)
```

```
In [51]: pre=rf.predict(x2_test)
```

```
In [52]: pre
```

```
Out[52]: array([42, 17, 17, ..., 75, 49, 54], dtype=int64)
```

```
In [53]: from sklearn.metrics import accuracy_score
```

```
In [54]: c=accuracy_score(y_test,pre)
```

```
In [55]: c
```

```
Out[55]: 0.9996419503938546
```

knn(k nearest neighbour)

```
In [56]: df.head()
```

```
Out[56]:
```

	Commodity	country	HSCode	value	year
0	71	0	5	0.000000	2018
1	22	0	7	12.380000	2018
2	21	0	8	268.600000	2018
3	16	0	9	35.480000	2018
4	72	0	11	63.289855	2018

```
In [58]: x3=df.iloc[:,[0,1,2,3,4]].values
          y=df.iloc[:,2].values
```

```
In [59]: from sklearn.model_selection import train_test_split
```

```
In [60]: x3_train, x3_test, y_train, y_test = train_test_split(x3, y, test_size=0.33, r
         : random_state=0)
```

```
In [61]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [62]: knn=KNeighborsClassifier()
```

```
In [63]: knn.fit(x3_train,y_train)
```

```
Out[63]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='minkowski',
         : metric_params=None, n_jobs=None, n_neighbors=5, p=2,
         : weights='uniform')
```

```
In [64]: pre=knn.predict(x3_test)
```

```
In [65]: pre
```

```
Out[65]: array([42, 26, 17, ..., 75, 49, 54], dtype=int64)
```

```
In [66]: from sklearn.metrics import accuracy_score
```

```
In [67]: d=accuracy_score(y_test,pre)
```

```
In [68]: d
```

```
Out[68]: 0.9014061584532257
```

svm(support vector machine)

```
In [69]: df.head()
```

```
Out[69]:
```

	Commodity	country	HSCode	value	year
0	71	0	5	0.000000	2018
1	22	0	7	12.380000	2018
2	21	0	8	268.600000	2018
3	16	0	9	35.480000	2018
4	72	0	11	63.289855	2018

```
In [70]: x4=df.iloc[:,[0,1,2,3,4]].values
         : y=df.iloc[:,2].values
```

```
In [71]: from sklearn.model_selection import train_test_split
```

```
In [72]: x4_train, x4_test, y_train, y_test = train_test_split(x4, y, test_size=0.33, r
        random_state=0)
```

```
In [73]: from sklearn.svm import LinearSVC
```

```
In [74]: svm=LinearSVC()
```

```
In [75]: svm.fit(x4_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\svm\base.py:931: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)

```
Out[75]: LinearSVC(C=1.0, class_weight=None, dual=True, fit_intercept=True,
        intercept_scaling=1, loss='squared_hinge', max_iter=1000,
        multi_class='ovr', penalty='l2', random_state=None, tol=0.0001,
        verbose=0)
```

```
In [76]: pre=svm.predict(x4_test)
```

```
In [77]: pre
```

```
Out[77]: array([54, 54, 54, ..., 54, 54, 54], dtype=int64)
```

```
In [78]: from sklearn.metrics import accuracy_score
```

```
In [79]: e=accuracy_score(y_test,pre)
```

```
In [80]: e
```

```
Out[80]: 0.019009179089903003
```

naive baiyes(nb)

```
In [81]: df.head()
```

```
Out[81]:
```

	Commodity	country	HSCode	value	year
0	71	0	5	0.000000	2018
1	22	0	7	12.380000	2018
2	21	0	8	268.600000	2018
3	16	0	9	35.480000	2018
4	72	0	11	63.289855	2018

```
In [82]: x5=df.iloc[:,[0,1,2,3,4]].values
        y=df.iloc[:,2].values
```



```
In [83]: from sklearn.model_selection import train_test_split
```

```
In [84]: x5_train, x5_test, y_train, y_test = train_test_split(x5, y, test_size=0.33, r
         random_state=0)
```

```
In [85]: from sklearn.naive_bayes import GaussianNB
```

```
In [86]: nb=GaussianNB()
```

```
In [87]: nb.fit(x5_train,y_train)
```

```
Out[87]: GaussianNB(priors=None, var_smoothing=1e-09)
```

```
In [92]: pre=nb.predict(x5_test)
```

```
In [93]: pre
```

```
Out[93]: array([42, 17, 17, ..., 75, 49, 54], dtype=int64)
```

```
In [94]: from sklearn.metrics import accuracy_score
```

```
In [95]: f=accuracy_score(y_test,pre)
```

```
In [96]: f
```

```
Out[96]: 1.0
```

```
In [ ]:
```